

The Jobber GraphQL API: A Comprehensive Technical Analysis and Developer's Guide

1. Executive Summary

This report provides a detailed technical analysis of the Jobber GraphQL API, its structure, capabilities, and underlying constraints. The analysis reveals that the API is built on a modern, single-endpoint GraphQL architecture, providing a robust and flexible foundation for custom integrations. Core capabilities include comprehensive data retrieval via queries, data modification via purpose-built mutations, and an event-driven system powered by webhooks. The platform's commitment to stability is demonstrated through a dual-layered, point-based rate-limiting system and a predictable, date-based versioning policy. While the API offers significant power, developers must adopt specific best practices—such as implementing efficient pagination, asynchronous webhook processing, and proactive error handling for throttling—to build resilient and scalable applications. The examination of community discussions reveals a discrepancy between the API's technical potential and the limitations of certain off-the-shelf integration tools, underscoring the value of direct, custom development for complex business logic.

2. Architectural Foundation: The Jobber GraphQL Endpoint & Developer Tooling

2.1 The Single Endpoint Model

The Jobber API is founded on a standard GraphQL architecture, centralizing all API interactions through a single, well-defined endpoint. All requests must be made via an HTTP POST to the URL <https://api.getjobber.com/api/graphql>.¹ As of April 2, 2024, the API strictly requires the

application/json content type, deprecating legacy formats such as application/x-www-form-urlencoded and multipart/form-data.¹ This shift towards a single, consistent content type streamlines the request process and aligns with modern API design principles. The choice of a single GraphQL endpoint, as opposed to a RESTful collection of endpoints, signifies a deliberate design philosophy focused on developer flexibility. Developers are empowered to retrieve or modify precisely the data they need in a single request, which reduces the number of calls and improves performance. This is a critical feature, especially when considering the API's query cost-based rate limiting, as it allows for more efficient data consumption.

2.2 Authentication and Authorization Framework (OAuth 2.0)

Access to the Jobber API is managed through the industry-standard OAuth 2.0 authorization framework.² This delegated authorization model ensures that a third-party app, referred to as the

Client, can access a user's Jobber account without ever handling their login credentials.⁴ The authorization flow involves the user, who is the

Resource Owner, first granting an authorization code, which the app then exchanges for a secure access token and refresh token.⁴ All subsequent GraphQL requests must include the

access token in an Authorization header, prepended with Bearer .¹ Scopes, defined during app creation, are crucial for controlling the specific data an application can read or write and for ensuring a streamlined OAuth screen experience for the user.⁵

An interesting aspect of the clientCreate mutation is the automatic capture of the app's name in the Lead source field for any client created after September 16, 2024.¹ This functionality is a strategic component that links a technical API capability to Jobber's business model. When an app on the marketplace creates a new client, Jobber can track and attribute that new business to the app. This provides a measurable and tangible value proposition for developers, allowing them to demonstrate their contribution to a client's growth. For Jobber, it creates a powerful incentive for developers to build applications that drive new client acquisition, solidifying the value of their developer ecosystem.⁶ This mutual benefit points to a

sophisticated and well-thought-out platform strategy.

2.3 The GraphQL Interface

The Jobber platform provides a live, interactive GraphQL environment within the Developer Center.⁶ This serves as a primary tool for developers to explore the API's schema, test queries and mutations, and confirm their functionality before writing production code.⁸ The live nature of this playground means that any queries or mutations executed directly affect the connected Jobber account, making it a powerful but potentially impactful testing tool.⁷ The provision of a full-featured GraphQL interface is a hallmark of a mature, developer-centric API. It significantly lowers the barrier to entry for developers and facilitates rapid prototyping and debugging, a practice Jobber explicitly recommends.⁹

2.4 The GraphQL Schema Structure

The GraphQL schema acts as a contract between the server and the client, defining all the capabilities of the API.²¹ Developers can view the full, up-to-date schema within the GraphQL tool.⁸

The schema is a collection of object types that contain fields.²¹ Types are declared using the type keyword, and fields are defined by their name and type.²¹ The schema uses specific formatting to indicate data constraints:

- A non-nullable field, which the API promises will always return a value, is marked with an exclamation mark !.²¹
- A list of items is indicated with square brackets [].²¹

The schema's entry points are the root operation types: Query for read operations and Mutation for write operations.²²

An example of an object type is Client, which represents a customer.⁸ Its schema defines the available fields and their types, such as:

- id: EncodedId! (a unique, non-null identifier)⁸
- balance: Float! (a non-null floating-point number)⁸
- companyName: String (a nullable string)⁸

- invoices: InvoiceConnection! (a non-null connection to a list of related invoices)⁸
- jobs: JobConnection! (a non-null connection to a list of related jobs)⁸
- requests: RequestConnection! (a non-null connection to a list of related requests)⁸

For modifying data, mutations often use input types to pass structured data.²³ These are declared with the

input keyword and serve as a way to group multiple input parameters.²³ By convention, input types are often named with

Input on the end (e.g., MessageInput).²³ Importantly,

input types can only contain scalar types, list types, and other input types, but cannot contain fields that are other object types.²³

3. Core API Capabilities: Data Interaction and Event Handling

3.1 GraphQL Queries for Data Retrieval

Jobber's GraphQL API uses queries to retrieve data. The schema includes a wide range of objects crucial to a home service business, such as Clients, Requests, Quotes, Jobs, and Invoices.⁸ A standard query structure, as exemplified by the

SampleQuery for clients, includes a top-level field that returns a connection object containing nodes (the actual data records) and totalCount (the total number of records).¹ This connection-based model strongly indicates the use of a cursor-based pagination system, a standard practice for GraphQL APIs dealing with large datasets.¹⁰

The query SampleQuery { clients { nodes { ... } totalCount } } structure¹ precisely matches the Relay Connection Specification, which is a key design pattern in GraphQL. The presence of

nodes and totalCount implies the existence of other fields like edges and pageInfo that are standard to this model. A cursor-based approach offers greater stability and reliability for paginating through dynamically changing data sets compared to a simple offset-based

model.¹¹ This model prevents records from being skipped or duplicated during pagination if new data is added or removed, a critical feature for a real-time system like Jobber. Therefore, the simple

nodes and totalCount fields are a subtle but powerful indicator of the API's technical maturity and robustness.

The relationships between key API objects are foundational to building effective integrations. The following table provides a high-level overview of these connections.

Table 1: Key API Objects and Relationships

Object	Related Objects	Description
Client	jobs, invoices, quotes, requests, properties	A central entity representing a customer. Most other objects are linked to a client. ⁸
Requests	client, jobs, quotes	Forms a client can fill to request work. Can be linked to a client and converted to jobs or quotes. ⁸
Quotes	client, jobs, requests, invoices	Proposals for work to be done. Can be converted to jobs. ⁸
Jobs	client, requests, quotes, invoices	The primary work object that represents a task to be completed. ⁸
Invoices	client, jobs, quotes	Financial records for services rendered. Can be linked to a client, job, or quote. ⁸

3.2 GraphQL Mutations for Data Modification

To modify data, developers use GraphQL mutations. These are designed to perform specific, purpose-built tasks, such as creating a new client via the `clientCreate` mutation.¹ Mutations accept

input objects to define the data to be changed and often return a payload with the modified object and a `userErrors` field for detailed error handling.¹ The API's schema defines a comprehensive set of mutations for managing core objects like clients, jobs, and invoices.¹³ For example, the

`clientCreate` mutation requires an input object containing fields like `firstName`, `lastName`, and `companyName` to create a new client.¹ This structured approach ensures data integrity and provides clear feedback to the developer about the status of the operation.

3.3 Event-Driven Integration with Webhooks

Jobber provides a webhook system to enable real-time, event-driven integrations.¹⁴ Webhooks are configured at the application level and trigger a

POST request to a developer-provided URL when a specific event occurs, such as `CLIENT_CREATE` or `CLIENT_UPDATE`.¹⁴ The platform enforces strict requirements on webhook processing, including a 1-second response time limit and a commitment to

at-least-once delivery.¹⁴ The

at-least-once delivery policy for webhooks means a developer's server might receive the same webhook payload multiple times. This necessitates building an idempotent system that can safely process duplicate payloads without causing side effects. For instance, a unique ID within the webhook payload can be used to check if a task has already been completed.

Furthermore, the 1-second response time limit for the initial POST request is a very tight constraint. A developer cannot perform a complex, time-consuming operation synchronously in response to the webhook. The appropriate strategy is to immediately return a 2xx status code and then queue the payload for asynchronous background processing.¹⁴ Failing to adhere to these best practices could lead to the app's webhooks being disabled, causing a critical failure in the integration.¹⁴ This design principle means a simple webhook implementation is insufficient; a robust system requires an asynchronous processing queue

and idempotent logic.

4. Navigating API Constraints: Performance, Pagination, and Rate Limiting

4.1 Dual-Layered Rate Limiting

Jobber's API employs two distinct rate limiters to ensure platform stability.¹⁵ The first is a DDoS protection middleware, which applies a hard limit of 2500 requests per 5 minutes on a per-app/account basis. Exceeding this limit results in a "429 Too Many Requests" error.¹⁵ This is a broad, first-line defense. The second, more sophisticated, system is the GraphQL query cost calculation. Every app/account combination has a pool of

currentlyAvailable points, which are consumed by each query and restored over time by a restoreRate.¹⁵ This mechanism is based on the leaky bucket algorithm, which provides a flexible and fair way to manage API consumption across the ecosystem.¹⁵

4.2 The Leaky Bucket Algorithm & Response-Based Throttling Feedback

The extensions.cost object in the API response is crucial for navigating the point-based rate limit.¹⁵ It provides real-time feedback on query cost and throttling status. The

requestedQueryCost is the estimated cost of the query, and if it exceeds the currentlyAvailable points, the API returns a THROTTLED error before the query is executed.¹⁵ The

actualQueryCost is the final cost subtracted from the currentlyAvailable points.¹⁵

This dual-layered rate-limiting system, particularly the extensions.cost object, shifts the burden of managing API consumption from a simple trial-and-error model to a data-driven, strategic approach. While many rate limits are a black box, Jobber's system provides a clear,

real-time feedback loop. A developer's application can parse the `currentlyAvailable` and `restoreRate` fields to make intelligent decisions. For instance, if a query is too expensive, the application can dynamically reduce the number of fields or items requested to stay within the limit. If the available points are low, the application can programmatically introduce a delay to allow points to accumulate, preventing a throttled response.¹⁵ This feedback mechanism allows for the creation of self-regulating, resilient applications that can adapt to API traffic without manual intervention.

The following table provides a breakdown of the key parameters within the `extensions.cost` object:

Table 2: GraphQL Query Cost Parameters

Parameter	Description
<code>requestedQueryCost</code>	The expected cost of the query. Used to determine if a query should be throttled.
<code>actualQueryCost</code>	The final amount of points consumed after the query has been resolved.
<code>maximumAvailable</code>	The maximum number of points that will ever be available for querying.
<code>currentlyAvailable</code>	The current number of points available to query with.
<code>restoreRate</code>	The rate at which points are added back to the <code>currentlyAvailable</code> amount per second.

4.3 Pagination and Filtering

The provided query examples indicate that Jobber's API uses a connection-based model, which is foundational to a cursor-based pagination system.¹ Jobber's documentation on rate

limits further mandates the use of

first or last arguments for connections to avoid a high requestedQueryCost.¹⁵ Failure to provide these arguments assumes the maximum number of records will be returned, leading to a higher cost and a greater likelihood of throttling.¹⁵ The design of the pagination model is inextricably linked to the query cost rate-limiting policy. Jobber's system incentivizes efficient data fetching. A developer who requests a large list of clients without pagination will incur a massive

requestedQueryCost, likely triggering a throttled response.¹⁵ This design choice forces developers to adopt the recommended practice of providing

first or last arguments, which in turn ensures that the API is used efficiently. The cost field is not merely a cap; it is a mechanism to enforce responsible API consumption and maintain platform stability for all users.

5. API Lifecycle Management: Versioning and Deprecation

5.1 Date-Based Versioning & Specification

Jobber's API uses a date-based versioning format (YYYY-MM-DD) for managing breaking changes.¹⁸ To specify a version, every request must include the

X-JOBBER-GRAPHQL-VERSION HTTP header.¹⁸ This explicit header requirement prevents unintended upgrades and provides developers with precise control over their integration's stability.¹⁸ The response includes version information under the

extensions key, providing real-time feedback on the version being used.¹⁸

The following table summarizes the critical aspects of Jobber's versioning policy:

Table 3: Jobber API Versioning Policy

Policy Component	Details
Version Format	YYYY-MM-DD
Required Header	X-JOBBER-GRAPHQL-VERSION
Minimum Support Period	12 months from the release of a newer version.
Maximum Accessibility	18 months from the version's release date.
Deprecation Warning	A warning is returned when a version will be unsupported in the next 3 months.
Action on Removal	The request is automatically upgraded to the next supported version.

5.2 Deprecation Policy and Changelog

Jobber's deprecation policy provides a predictable path for developers. Old versions are supported for a minimum of 12 months, with a maximum accessibility of 18 months from their release date.¹⁷ When a version becomes unsupported, the API returns a warning message, and upon removal, it automatically upgrades the request to the next supported version.¹⁸ All breaking changes are documented in a publicly available changelog.¹⁹

While the deprecation policy is well-defined, the automatic upgrade to the oldest supported version after removal, without prior warning, introduces a potential risk for integration stability. The policy states that once a version is removed, any request using it will be automatically upgraded to the "next supported version (oldest supported version)".¹⁸ The documentation also notes that this upgrade can occur "at any time" without warning after the version becomes unsupported. For an application that was designed for a specific schema, a silent, unexpected upgrade could lead to runtime errors as the application might be receiving data in a different format than it expects. This places a strong emphasis on the developer's responsibility to proactively monitor deprecation warnings and the changelog, as relying on

the "automatic upgrade" could lead to a critical failure of the integration.

6. Real-World Developer Insights and Strategic Recommendations

6.1 Community-Sourced Challenges and Workarounds

Community discussions reveal that while the API is powerful, developers face challenges with specific integration scenarios.²⁰ Issues include limitations with off-the-shelf tools like Zapier, which may not pass all necessary data fields, such as custom fields, or can quickly exhaust rate limits.²⁰ This often necessitates the use of more flexible but complex tools like webhooks or custom-developed solutions to achieve the desired automations.²⁰ For example, a developer had to hire a third-party developer to create a custom Google Function to automate the creation of a quote from a CompanyCam project.²⁰ Inconsistent report formats and the inability to automate certain reports were also cited as significant pain points, requiring manual workarounds to update business intelligence dashboards.²⁰

6.2 Best Practices for Building Resilient Integrations

Based on this analysis, building a successful Jobber integration requires a strategic approach that goes beyond simple API calls.

- **Embrace Asynchronous Processing:** An asynchronous queue should be used to process webhook payloads to adhere to the 1-second response time limit and prevent the app's webhooks from being disabled.¹⁴
- **Implement Idempotency:** The at-least-once delivery policy for webhooks requires building idempotent logic to prevent duplicate processing.¹⁴
- **Strategically Manage Query Cost:** Developers should actively monitor the `extensions.cost` object to prevent throttling and introduce dynamic delays or reduce query complexity when necessary.¹⁵
- **Use first or last Arguments:** It is a best practice to always include these arguments in connection queries to manage costs and avoid high `requestedQueryCost` values.¹⁵

- **Proactively Monitor Versioning:** Developers must keep up to date with the changelog and deprecation notices to ensure the application remains compatible with the API's latest version.¹⁸
-

7. Conclusion: Final Assessment and Strategic Recommendations

The Jobber GraphQL API is a technically robust and developer-friendly platform that provides the necessary tools for building powerful integrations. Its single-endpoint, query-based architecture and mature authentication framework simplify data access, while its sophisticated rate-limiting and versioning policies provide the necessary guardrails for stability. The platform's investment in developer tooling, particularly GraphiQL, and event-driven webhooks, signals a clear commitment to enabling a rich ecosystem. However, developers must be prepared to navigate nuanced challenges, particularly the tight constraints on webhook processing and the need to strategically manage query costs. For organizations seeking to build complex, business-critical automations, a direct-to-API development approach, rather than relying solely on pre-built third-party connectors, is the recommended path to fully leverage the platform's capabilities and overcome the limitations cited in the developer community. This strategy will allow for the implementation of custom logic that can handle the unique challenges of data synchronization, custom fields, and real-time business processes, ensuring a more resilient and effective integration.

Works cited

1. API Queries and Mutations - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/using_jobbers_api/api_queries_and_mutations/
2. App Listing Details - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/publishing_your_app/app_listing_details/
3. Custom Integrations - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/custom_integrations/
4. App Authorization (OAuth 2.0) - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/building_your_app/app_authorization
5. Getting Started - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/getting_started/
6. Jobber's Developer Center, accessed August 17, 2025, <https://developer.getjobber.com/>
7. Developer Center - Jobber Help Center, accessed August 17, 2025, <https://help.getjobber.com/hc/en-us/articles/25924078048151-Developer-Center>

8. Jobber's API - Jobber's Developer Center, accessed August 17, 2025, <https://developer.getjobber.com/docs/>
9. Testing Your Application - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/building_your_app/testing_your_app/
10. Pagination - GraphQL, accessed August 17, 2025, <https://graphql.org/learn/pagination/>
11. Cursor based pagination - Medium, accessed August 17, 2025, <https://medium.com/@nimmikrishnab/cursor-based-pagination-37f5fae9f482>
12. Cursor pagination: how it works and its pros and cons - Merge.dev, accessed August 17, 2025, <https://www.merge.dev/blog/cursor-pagination>
13. How to build a Jobber API integration - Rollout, accessed August 17, 2025, <https://rollout.com/integration-guides/jobber/sdk/step-by-step-guide-to-building-a-jobber-api-integration-in-java>
14. Setting up Webhooks - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/using_jobbers_api/setting_up_webhooks
15. API Rate Limits - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/using_jobbers_api/api_rate_limits/
16. API Rate Limits - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/using_jobbers_api/api_rate_limits
17. Building an App in Jobber Platform - DEV Community, accessed August 17, 2025, <https://dev.to/jobber/building-an-app-in-jobber-platform-5259>
18. API Versioning - Jobber's Developer Center, accessed August 17, 2025, https://developer.getjobber.com/docs/using_jobbers_api/api_versioning
19. Changelog - Jobber's Developer Center, accessed August 17, 2025, <https://developer.getjobber.com/docs/changelog/>
20. Best Jobber Automations | The Home Service Community - 208, accessed August 17, 2025, <https://community.getjobber.com/discussions/operations-forum/best-jobber-automations/208/replies/703>
21. Schema definition language (SDL) - Apollo GraphQL, accessed August 17, 2025, <https://www.apollographql.com/tutorials/lift-off-part1/03-schema-definition-language-sdl>
22. Schemas and Types - GraphQL, accessed August 17, 2025, <https://graphql.org/learn/schema/>
23. Mutations and Input Types - GraphQL.js, accessed August 17, 2025, <https://www.graphql-js.org/docs/mutations-and-input-types/>